

ECE408 Spring 2020

Applied Parallel Programming

Lecture 16

Parallel Computation Patterns –  
Parallel Scan (Prefix Sum)

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

1

1

Objective

- to learn parallel scan (prefix sum) algorithms based on reductions and reverse reductions
- to learn the concept of double buffering
- to understand tradeoffs between work efficiency and latency
- to learn how to develop hierarchical algorithms (across multiple kernels)

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

2

2

Scan Includes all Partial Results

Reductions are a simplified form of scans.

In scan / parallel prefix,

- we need all of the partial sums
- (or whatever the operator might be).

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

3

3

(Inclusive) Scan (Prefix-Sum) Definition

**Definition:** The scan operation takes a binary associative operator  $\oplus$ , and an array of  $n$  elements  $[x_0, x_1, \dots, x_{n-1}]$ , and returns the prefix-sum array

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

**Example:** If  $\oplus$  is addition, the scan operation on the array  $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$ , returns  $[3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22 \ 25]$ .

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

4

4

## Example: Sharing a Big Sandwich

You order a 100-inch sandwich to feed 10 people, and you know how much each person wants in inches:

[3 5 2 7 28 4 3 0 8 1].

**How do you cut the bread quickly?**

**How much of the sandwich is left over?**

Method 1: sequentially!

Cut 3 inches, then cut 5 inches, then ...

Method 2: **calculate cutting offsets with prefix-sum**

[3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

5

5

## Typical Applications of Scan

A simple and useful parallel building block.

Convert sequential recurrences

```
for(j=1;j<n;j++)  
    out[j] = out[j-1] + f(j);
```

into parallel:

```
forall(j) { temp[j] = f(j) };  
scan(out, temp);
```

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

6

6

## Typical Applications of Scan

- Useful for many parallel algorithms:
  - radix sort
  - quicksort
  - String comparison
  - Lexical analysis
  - Stream compaction
  - Polynomial evaluation
  - Solving recurrences
  - Tree operations
  - Histograms
  - Etc.

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

7

7

## Other Applications

- Assigning camp slots
- Assigning farmer market space
- Allocating memory to parallel threads
- Allocating memory buffer to communication channels
- ...

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

8

8

## An Inclusive Sequential Scan

Given a sequence  $[x_0, x_1, x_2, \dots]$

Calculate output  $[y_0, y_1, y_2, \dots]$

Such that

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

...

*Using a recursive definition*

$$y_i = y_{i-1} + x_i$$

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

9

9

## An Sequential C Implementation

```
y[0] = x[0];  
for (i = 1; i < Max_i; i++) y[i] = y[i-1] + x[i];
```

Computationally efficient:

N additions needed for N elements -  $O(N)$ !

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

10

10

## A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each y element
- Have every thread to add up all x elements needed for the y element

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

“Parallel programming is easy as long as you do not care about performance.”

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

11

11

## Parallel Inclusive Scan using Reduction Trees

Calculate each output element as the reduction of all previous elements

- Some reduction partial sums will be shared among the calculation of output elements
- Based on hardware added design by Peter Kogge and Harold Stone at IBM in the 1970s – Kogge-Stone Trees
- Goal: low latency

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

12

12

## A Kogge-Stone Parallel Scan Algorithm

T	3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---	---

1. Load input from global memory into shared memory array T


Each thread loads one value from the input (global memory) array into shared memory array T.

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

13

13

## A Kogge-Stone Parallel Scan Algorithm



1. (previous slide)
2. Assuming  $n$  is a power of 2. Iterate  $\log(n)$  times, stride from 1 to  $n/2$ . Threads  $stride$  to  $n-1$  active: add pairs of elements that are  $stride$  elements apart.

- Active threads:  $stride$  to  $n-1$  ( $n - stride$  active threads)
- Thread  $j$  adds elements  $T[j]$  and  $T[j - stride]$  and writes result into element  $T[j]$
- Each iteration requires two syncthreads
  - make sure that input is in place
  - make sure that all input elements have been used


Iteration #1  
Stride = 1

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

14

14

## A Kogge-Stone Parallel Scan Algorithm



1. (previous slide)
2. Assuming  $n$  is a power of 2. Iterate  $\log(n)$  times, stride from 1 to  $n/2$ . Threads  $stride$  to  $n-1$  active: add pairs of elements that are  $stride$  elements apart.

- Active threads:  $stride$  to  $n-1$  ( $n - stride$  active threads)
- Thread  $j$  adds elements  $T[j]$  and  $T[j - stride]$  and writes result into element  $T[j]$
- Each iteration requires two syncthreads
  - `syncthreads();` // make sure that input is in place
  - `float temp = T[j] + T[j - stride];`
  - `syncthreads();` // make sure that previous output has been consumed
  - `T[j] = temp;`


Iteration #1  
Stride = 1

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

15

15

## A Kogge-Stone Parallel Scan Algorithm




Iteration #2  
Stride = 2

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

16

16

## A Kogge-Stone Parallel Scan Algorithm




Iteration #3  
Stride = 4

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

17

17

## Sharing Computation in Kogge-Stone



Iteration #3  
Stride = 4

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

18

18

## Double Buffering


- Use two copies of data T0 and T1
- Start by using T0 as input and T1 as output
- Switch input/output roles after each iteration
  - Iteration 0: T0 as input and T1 as output
  - Iteration 1: T1 as input and T0 as output
  - Iteration 2: T0 as input and T1 as output
- This is typically implemented with two pointers, source and destination that swap their contents from one iteration to the next
- This eliminates the need for the second `__syncthreads()` call

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

19

19

## A Double-Buffered Kogge-Stone Parallel Scan Algorithm




Iteration #1  
Stride = 1

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

20

20

## A Kogge-Stone Parallel Scan Algorithm




Iteration #2  
Stride = 2

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

21

21

## Sharing Computation in Kogge-Stone



- Each iteration requires only one syncthreads()
  - syncthreads(); // make sure that input is in place
  - float destination[j] = source[j] + source[j-stride];
  - temp = destination; destination = source; source = temp;
- After the loop, write destination contents to global memory

Iteration #3  
Stride = 4

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

22

22

## Work Efficiency Analysis


- A Kogge-Stone scan kernel executes  $\log(n)$  parallel iterations
  - The steps do  $(n-1)$ ,  $(n-2)$ ,  $(n-4)$ , ...,  $(n-n/2)$  add operations each
  - Total # of add operations:  $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$  work
- This scan algorithm is not very work efficient
  - Sequential scan algorithm does  $n$  adds
  - A factor of  $\log(n)$  hurts: 20x for 1,000,000 elements!
  - Typically used within each block, where  $n \leq 1,024$
- A parallel algorithm can be slow when execution resources are saturated due to low work efficiency

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

23

23

## A Kogge-Stone Parallel Scan Algorithm



© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

24

24

## Improving Efficiency

- A common parallel algorithm pattern:
  - Balanced Trees*
  - Build a balanced binary tree on the input data and sweep it to and from the root
  - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
  - Traverse down from leaves to root building partial sums at internal nodes in the tree
    - Root holds sum of all leaves
  - Traverse back up the tree building the scan from the partial sums


© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

25

25

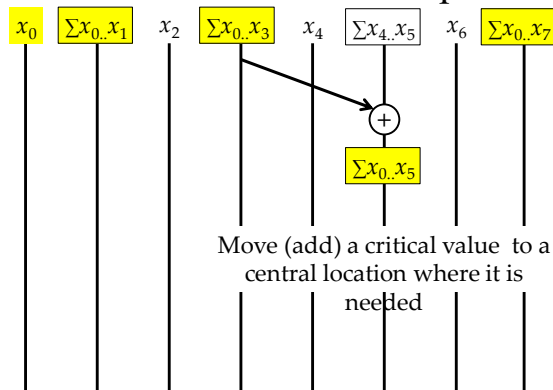
## Brent-Kung Parallel Scan

### - Reduction Step



26

## Inclusive Post-Scan Step

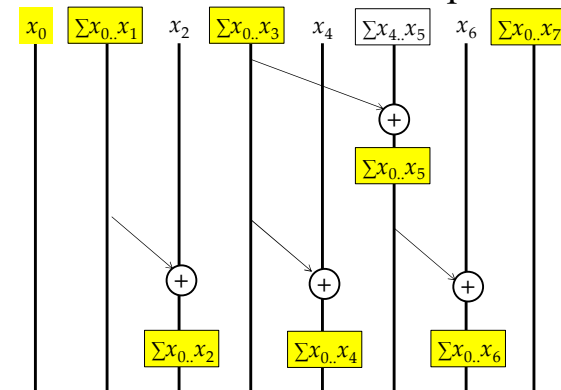


© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

27

27

## Inclusive Post Scan Step




© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

28

28

## Putting it Together (Data View)



29

## Reduction Step Kernel Code

```
// float T[2*BLOCK_SIZE] is in shared memory
// for previous slide, BLOCK_SIZE is 8
int stride = 1;
while(stride < 2*BLOCK_SIZE)
{
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE && (index-stride) >= 0)
        T[index] += T[index-stride];
    stride = stride*2;
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

30

30

## Reduction Step Kernel Code

// float T[2\*BLOCK\_SIZE] is in shared memory

```
int stride = 1;
while(stride < 2*BLOCK_SIZE)
{
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE && (index-stride) >= 0)
        T[index] += T[index-stride];
    stride = stride*2;
    __syncthreads();
}
```


// For previous example,  
// threadIdx.x+1 = 1, 2, 3, 4, 5, 6, 7, 8  
// stride = 1, index = 1, 3, 5, 7, 9, 11, 13, 15

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

31

31

## Putting it Together



© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

32



## Post Scan Step

```

int stride = BLOCK_SIZE/2;
while(stride > 0)
{
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if((index+stride) < 2*BLOCK_SIZE)
    {
        T[index+stride] += T[index];
    }
    stride = stride / 2;
}

```

// for the previous example,  
 // BLOCK\_SIZE is 8  
 // stride will go 4, 2, 1  
 // for the first iteration, the active thread  
 // will be thread 0, with index = 7 and  
 // index+stride = 11

© David Kirk/NVIDIA and Wen-mei W. Hwu  
 ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

33

33

## Work Analysis


- The parallel Inclusive Scan executes  $2 * \log(n)$  parallel iterations
  - $\log(n)$  in reduction and  $\log(n)$  in post scan
  - The iterations do  $n/2, n/4, \dots, (2-1), \dots, (n/4-1), (n/2-1)$  useful adds
  - In our example,  $n = 16$ , the number of useful adds is  $16/2 + 16/4 + 16/8 + 16/16 + (16/8-1) + (16/4-1) + (16/2-1)$
  - Total adds:  $(n-1) + (n-2) - (\log(n) - 1) = 2*(n-1) - \log(n) \rightarrow O(n)$  work
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
  - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware

© David Kirk/NVIDIA and Wen-mei W. Hwu  
 ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

34

34

## Kogge-Stone vs. Brent-Kung



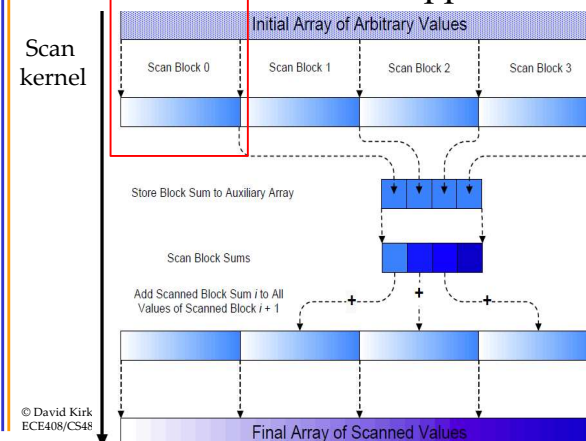
- Brent-Kung uses half the number of threads compared to Kogge-Stone
  - Each thread should load two elements into the shared memory
- Brent-Kung takes twice the number of steps compared to Kogge-Stone
  - Kogge-Stone is more popular for parallel scan with blocks in GPUs

© David Kirk/NVIDIA and Wen-mei W. Hwu  
 ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

35

35

## Overall Flow of Complete Scan A Hierarchical Approach



© David Kirk  
 ECE408/CS48

36

36

## Using Global Memory Contents in CUDA

- Data in registers and shared memory of one thread block are not visible to other blocks
- To make data visible, the data has to be written into global memory
- However, any data written to the global memory are not visible until a memory fence. This is typically done by terminating the kernel execution
- Launch another kernel to continue the execution. The global memory writes done by the terminated kernels are visible to all thread blocks.

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

37

37

## Scan of Arbitrary Length Input

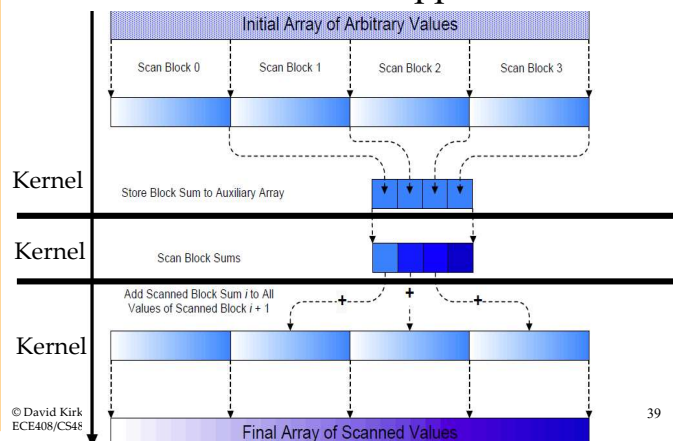
- Build on the scan kernel that handles up to  $2 \times \text{blockDim.x}$  elements from Brent-Kung
  - For Kogge-Stone, have each section of  $\text{blockDim.x}$  elements assigned to a block
- Have each block write the sum of its section into a Sum array using its  $\text{blockIdx.x}$  as index
- Run parallel scan on the Sum array
  - May need to break down Sum into multiple sections if it is too big for a block
- Add the scanned Sum array values to the elements of corresponding sections

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

38

38

## Overall Flow of Complete Scan A Hierarchical Approach



39

39

## (Exclusive) Scan Definition

**Definition:** The exclusive scan operation takes a binary associative operator  $\oplus$ , and an array of  $n$  elements  $[x_0, x_1, \dots, x_{n-1}]$

and returns the array

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

**Example:** If  $\oplus$  is addition, then the exclusive scan operation on  $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$ , would return  $[0 \ 3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22]$ .

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

40

40

## Why Exclusive Scan

- To find the beginning address of allocated buffers
- Inclusive and Exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

41

41

## A simple exclusive scan kernel

- Adapt an inclusive, Kogge-Stone scan kernel
  - Block 0:
    - Thread 0 loads 0 into (shared) XY[0]
    - Other threads load (global) X[threadIdx.x-1] into XY[threadIdx.x]
  - All other blocks:
    - All thread load X[blockIdx.x\*blockDim.x+threadIdx.x-1] into XY[threadIdx.x]
- Similar adaption for Brent-Kung kernel but pay attention that each thread loads two elements
  - Only one zero should be loaded
  - All elements should be shifted by only one position
- Intellectual contribution vs. practical contribution

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

42

42

## ANY MORE QUESTIONS?

© David Kirk/NVIDIA and Wen-mei W. Hwu  
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

43

43